



ITAI 1378 • COMPUTER VISION — FINAL PROJECT

PPE Compliance Checker

Automated detection of hard hats, hi-vis vests, and safety goggles — flagging every worker as compliant or non-compliant.

Trilok Kalani • Matthew Jenkins

Tier 2 • Final system trained and evaluated — mAP@50 0.61 • ~20 ms/image (T4)

GitHub: github.com/Tiskalani/ITAI-1378-Final_PPE-Compliance-Checker

The Problem

Preventable injuries + compliance liability on the job site

Workers who skip required PPE face some of the most common — and most preventable — worksite injuries: head trauma, struck-by incidents, and eye injuries.

Manual spot-checks don't scale. One safety officer cannot watch every worker continuously, and violations are logged inconsistently, if at all.

The result is both a safety gap and an OSHA-compliance liability.

1

Who cares

Safety officers / EHS managers, superintendents, construction & manufacturing firms, insurers

2

Why it matters

Head, eye & struck-by hazards are leading, largely preventable injuries

3


Cost of failure

Direct OSHA citations + injury and liability cost

The Solution

One sentence, one model

An image of a work area goes into a fine-tuned detector that locates each worker and each PPE item, then labels every person “Compliant” or “Non-compliant.”



V2

Next step (V2): extend to a live or recorded video stream — track each worker across frames and log compliance events with timestamps, making it a deployable, portfolio-ready safety-monitoring tool.

Technical Approach

Object detection, done with the fastest reliable tool

1

TECHNIQUE

Multi-class object detection

Locate + identify people and several gear types in one pass.

2

MODEL

YOLO11 (Ultralytics), fine-tuned

Best accuracy-per-training-hour; trains on free Colab GPUs.

3

FRAMEWORK

PyTorch + Ultralytics

Standard, free, Colab-ready; built-in tracking for V2 video.

4

COMPLIANCE LOGIC

Rule-based, per worker

Associate gear to person; require detected hat + vest (goggles configurable).

Why YOLO11: top accuracy-per-training-hour for multi-class detection on a modest budget, and its inference speed is what carries the same model into the real-time V2 capstone without a rewrite.

Data Used

Public, already-labeled datasets — no manual collection

Roboflow Universe — PPE

Thousands of labeled images (a common set: ~19K / 3.6K / 1.9K)
person, hardhat, vest, goggles + NO- variants

SH17

8,099 images · ~75,994 instances
17 PPE + body-part classes (V2 scale-up)

Ultralytics Construction-PPE ← used for V1

1,416 images (1,132 train / 143 val / 141 test)
11 classes incl. helmet, vest, goggles, no_helmet, no_goggle


V1 uses Construction-PPE

- 1,416 images total
- 1,132 train / 143 val / 141 test
- Auto-downloads (178 MB), zero setup
- Standardize 640×640, light augmentation

Size: ample for a first fine-tune; scale to SH17 (8,099) for V2.

System Diagram

Input → inference → compliance decision



V2

Capstone: video stream → track workers across frames (ByteTrack) → timestamped compliance log.

Success Metrics

Targets vs. what the final model achieved

≥ 0.85

mAP@50 (target)

Detection quality across all classes

0.61

mAP@50 (achieved)

40 epochs, 1,416 imgs, held-out val

0.30

mAP@50-95 (achieved)

Stricter localization metric

Supporting checks

- Precision 0.66 · Recall 0.59. Strong on people and common gear (mAP@50 0.80–0.91); weak on rare violation classes (no_goggle 0.23, no_boots 0.16).
- Latency: ~20 ms/image end-to-end on a T4 — the < 1 s/image target met; ≥ 15 FPS is the V2 video bar.
- 0.61 on a 1,416-image set is an honest baseline — the fix is more data (scale to SH17), not a fancier metric.

Final result YOLO11s · 40 epochs · Tesla T4 · mAP@50 0.61 · mAP@50-95 0.30 · P 0.66 · R 0.59 (curves + confusion matrix on the following slides)

How the Six Weeks Went

The plan vs. what actually happened

1 Set up Colab; downloaded and explored Construction-PPE **Done — 1,416 imgs in YOLO format**

2 Fine-tuned YOLO11s — 40 epochs, ~17 min on a T4 **Done — first real mAP: 0.61**

3 Added per-worker compliance logic **Done — end-to-end labels on images**

4 Evaluated on the held-out val set (143 imgs) **0.61 vs 0.85 target — gap traced to rare no_* classes**

5 Built the demo — 6 annotated sample outputs **Done — samples in results/**

6 Laptop lost — recovered all work from GitHub; docs and slides rebuilt **Version control saved the project**

Challenges & Lessons Learned

What actually happened — and what saved us

CHALLENGE

Rare violation classes score low — no_goggle 0.23, no_boots 0.16 mAP@50

Small dataset caps accuracy (mAP@50 0.61)

Midterm per-class table mislabeled mAP@50-95 as mAP@50

Laptop lost near the deadline — all local files gone

OUTCOME

Diagnosed: too few no_* examples. V2 fix: SH17 (8,099 imgs) or oversampling


Improvement path: more data + larger variant (s→m) + more epochs

Caught in final review — common gear is actually 0.80–0.91 mAP@50

Everything recovered from the public GitHub repo — commit early, commit often

System in Action — Sample Outputs

Real outputs of the trained model — src/predict_compliance.py




How to read

- Green box = worker judged COMPLIANT
- Red box = NON-COMPLIANT, missing gear listed
- Rule: vest on person, no no_helmet / no_goggle
- All 6 samples in results/

Training Results

YOLO11s • 40 epochs • Tesla T4



Read

- Losses fall smoothly
- mAP climbs, plateaus ~epoch 35-40
- No overfitting or instability
- **mAP@50 → 0.61**

Per-Class Results

Where the model is strong vs. weak

Per-class mAP@50

Person	0.907	none	0.552
vest	0.850	no_helmet	0.448
gloves	0.826	no_gloves	0.289
goggles	0.816	no_goggle	0.233
helmet	0.803	no_boots	0.164
boots	0.799		

Insight: strong on people and common gear (0.80–0.91 mAP@50); weak on the rare no_* violation classes the compliance rule depends on (0.16–0.45). Fix = more violation examples (SH17 / oversampling).

